

## V1 CNN:



First version of our CNN, testing only on logistic regression, SVM, and KNN. No data augmentation was added.


## MobileNetV2 results:

Logistic Regression Accuracy: 0.5214285714285715

SVM Accuracy: 0.4857142857142857

KNN Accuracy: 0.4857142857142857








## ResNet50 Results:

Logistic Regression Accuracy: 0.8214285714285714

SVM Accuracy: 0.7642857142857142

KNN Accuracy: 0.6785714285714286








## EfficientNetB0 results:

Logistic Regression Accuracy: 0.6785714285714286

SVM Accuracy: 0.5071428571428571

KNN Accuracy: 0.4785714285714286



**CNN V2 without data augmentation:**

Updated version of CNN now with classification report and misclassified images for fine tuned model.

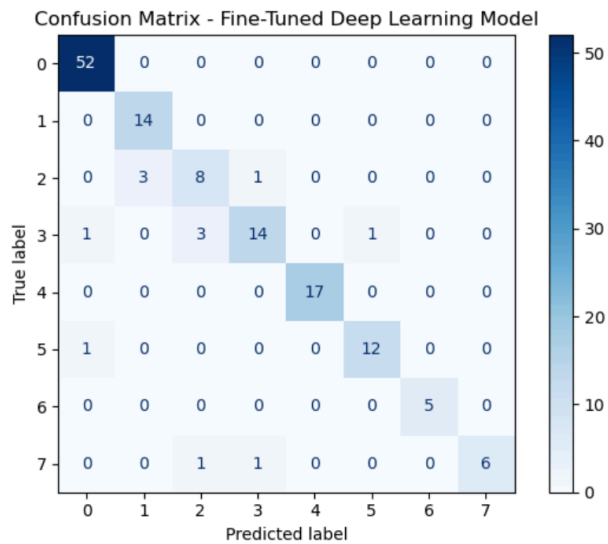
### ResNet50 results:

Fine-tuned Custom Model Accuracy: 0.9143

Logistic Regression Accuracy: 0.8286

SVM Accuracy: 0.8214

KNN Accuracy: 0.7429



### Classification Report - Fine-tuned Custom Model

	precision	recall	f1-score	support
0	0.96	1.00	0.98	52
1	0.82	1.00	0.90	14
2	0.67	0.67	0.67	12
3	0.88	0.74	0.80	19
4	1.00	1.00	1.00	17
5	0.92	0.92	0.92	13
6	1.00	1.00	1.00	5
7	1.00	0.75	0.86	8
accuracy			0.91	140
macro avg	0.91	0.88	0.89	140
weighted avg	0.92	0.91	0.91	140

Total misclassified images: 12




Misclassified Images - Fine-tuned Model



**EfficientNetB0 results:**

Fine-tuned Custom Model Accuracy: 0.8786  
Logistic Regression Accuracy: 0.8000  
SVM Accuracy: 0.7714  
KNN Accuracy: 0.7143



Classification Report - Fine-tuned Custom Model

	precision	recall	f1-score	support
0	0.95	1.00	0.97	52
1	0.86	0.86	0.86	14
2	0.58	0.58	0.58	12
3	0.76	0.68	0.72	19
4	0.94	1.00	0.97	17
5	1.00	0.85	0.92	13
6	1.00	0.80	0.89	5
7	0.78	0.88	0.82	8
accuracy			0.88	140
macro avg	0.86	0.83	0.84	140
weighted avg	0.88	0.88	0.88	140


Total misclassified images: 17

Misclassified Images - Fine-tuned Model

True: 0to25percentPartialEclipse  
Pred: 26to55percentPartialEclipse



True: 0to25percentPartialEclipse  
Pred: 26to55percentPartialEclipse



True: 26to55percentPartialEclipse  
Pred: 0to25percentPartialEclipse



True: 26to55percentPartialEclipse  
Pred: 56to95percentPartialEclipse




True: 26to55percentPartialEclipse  
Pred: 0to25percentPartialEclipse



True: 26to55percentPartialEclipse  
Pred: 56to95percentPartialEclipse



True: 26to55percentPartialEclipse  
Pred: 56to95percentPartialEclipse



True: 56to95percentPartialEclipse  
Pred: NotASolarEclipse



True: 56to95percentPartialEclipse  
Pred: TotalSolarEclipse



True: 56to95percentPartialEclipse  
Pred: 26to55percentPartialEclipse



True: 56to95percentPartialEclipse  
Pred: NotASolarEclipse



True: 56to95percentPartialEclipse  
Pred: 26to55percentPartialEclipse



True: 56to95percentPartialEclipse  
Pred: 26to55percentPartialEclipse



True: DiamondRing\_BaileysBeads\_SolarEclipse  
Pred: TotalSolarEclipse




True: DiamondRing\_BaileysBeads\_SolarEclipse  
Pred: TotalSolarEclipse



True: Flats  
Pred: Darks



True: NotASolarEclipse  
Pred: 56to95percentPartialEclipse





## CNN V2 with data augmentation:

Adding data augmentation to the CNN helped improve accuracy a bit and gave us our best results using ResNet50 with our custom model.


### ResNet50 Results:

Fine-tuned Custom Model Accuracy: 0.9571

Logistic Regression Accuracy: 0.8286

SVM Accuracy: 0.8214

KNN Accuracy: 0.7429




Classification Report - Fine-tuned Custom Model

	precision	recall	f1-score	support
0	1.00	1.00	1.00	52
1	1.00	0.93	0.96	14
2	0.79	0.92	0.85	12
3	0.89	0.84	0.86	19
4	1.00	1.00	1.00	17
5	0.93	1.00	0.96	13
6	1.00	1.00	1.00	5
7	1.00	0.88	0.93	8
accuracy			0.96	140
macro avg	0.95	0.95	0.95	140
weighted avg	0.96	0.96	0.96	140

Total misclassified images: 6

Misclassified Images - Fine-tuned Model

True: 0to25percentPartialEclipse  
Pred: 26to55percentPartialEclipse



True: 56to95percentPartialEclipse  
Pred: 26to55percentPartialEclipse




True: 26to55percentPartialEclipse  
Pred: 56to95percentPartialEclipse




True: 56to95percentPartialEclipse  
Pred: 26to55percentPartialEclipse



True: 56to95percentPartialEclipse  
Pred: DiamondRing\_BaileysBeads\_SolarEclipse



True: NotASolarEclipse  
Pred: 56to95percentPartialEclipse



## NOTE:

**The accuracy varied depending on the machine it was run on.**


*The following shows the accuracies running on a machine with the following specifications:*

- Processor: AMD Ryzen 9 5900X 12-Core Processor, 3701 MHz, 12 Core(s), 24 Logical Processors
- RAM: 32 GB
- Graphics Card: EVGA GeForce RTX 3080


Best accuracy: 92.86%

```
Fine-tuned Custom Model Accuracy: 0.9286  
Logistic Regression Accuracy: 0.8286  
SVM Accuracy: 0.8071  
KNN Accuracy: 0.7286
```


<Figure size 1000x800 with 0 Axes>




<Figure size 1000x800 with 0 Axes>



<Figure size 1000x800 with 0 Axes>



<Figure size 1000x800 with 0 Axes>



```

Classification Report - Fine-tuned Custom Model
precision    recall  f1-score   support

0           0.98     0.96     0.97     51
1           0.92     0.85     0.88     13
2           0.83     0.77     0.80     13
3           0.86     0.95     0.90     19
4           1.00     0.94     0.97     18
5           0.81     1.00     0.90     13
6           1.00     1.00     1.00      5
7           1.00     0.88     0.93      8

accuracy          0.93     140
macro avg         0.92     0.92     0.92     140
weighted avg      0.93     0.93     0.93     140

```

```

Classification Report - Logistic Regression
precision    recall  f1-score   support

0           0.89     1.00     0.94     51
1           0.69     0.69     0.69     13
2           0.47     0.54     0.50     13
3           0.69     0.58     0.63     19
4           1.00     1.00     1.00     18
5           1.00     0.69     0.82     13
6           0.83     1.00     0.91      5
7           1.00     0.75     0.86      8

accuracy          0.83     140
macro avg         0.82     0.78     0.79     140
weighted avg      0.84     0.83     0.83     140

```

```

Classification Report - SVM
precision    recall  f1-score   support

0           0.78     1.00     0.88     51
1           0.83     0.77     0.80     13
2           0.58     0.54     0.56     13
3           0.71     0.63     0.67     19
4           1.00     0.94     0.97     18
5           1.00     0.31     0.47     13
6           1.00     1.00     1.00      5
7           0.88     0.88     0.88      8

accuracy          0.81     140
macro avg         0.85     0.76     0.78     140
weighted avg      0.82     0.81     0.79     140

```

```

Classification Report - KNN
precision    recall  f1-score   support

0           0.75     1.00     0.86     51
1           0.55     0.85     0.67     13
2           0.36     0.31     0.33     13
3           0.69     0.47     0.56     19
4           0.94     0.94     0.94     18
5           1.00     0.15     0.27     13
6           1.00     1.00     1.00      5
7           1.00     0.38     0.55      8

accuracy          0.73     140
macro avg         0.79     0.64     0.65     140
weighted avg      0.76     0.73     0.69     140

```

Total misclassified images: 10

Misclassified Images - Fine-tuned Model

True: 26to55percentPartialEclipse  
Pred: 56to95percentPartialEclipse



True: 0to25percentPartialEclipse  
Pred: 26to55percentPartialEclipse




True: 0to25percentPartialEclipse  
Pred: 26to55percentPartialEclipse



True: 26to55percentPartialEclipse  
Pred: 0to25percentPartialEclipse



True: TotalSolarEclipse  
Pred: DiamondRing\_BaileysBeads\_SolarEclipse



True: 56to95percentPartialEclipse  
Pred: DiamondRing\_BaileysBeads\_SolarEclipse




True: Dark  
Pred: TotalSolarEclipse




True: 26to55percentPartialEclipse  
Pred: 56to95percentPartialEclipse



True: TotalSolarEclipse  
Pred: DiamondRing\_BaileysBeads\_SolarEclipse



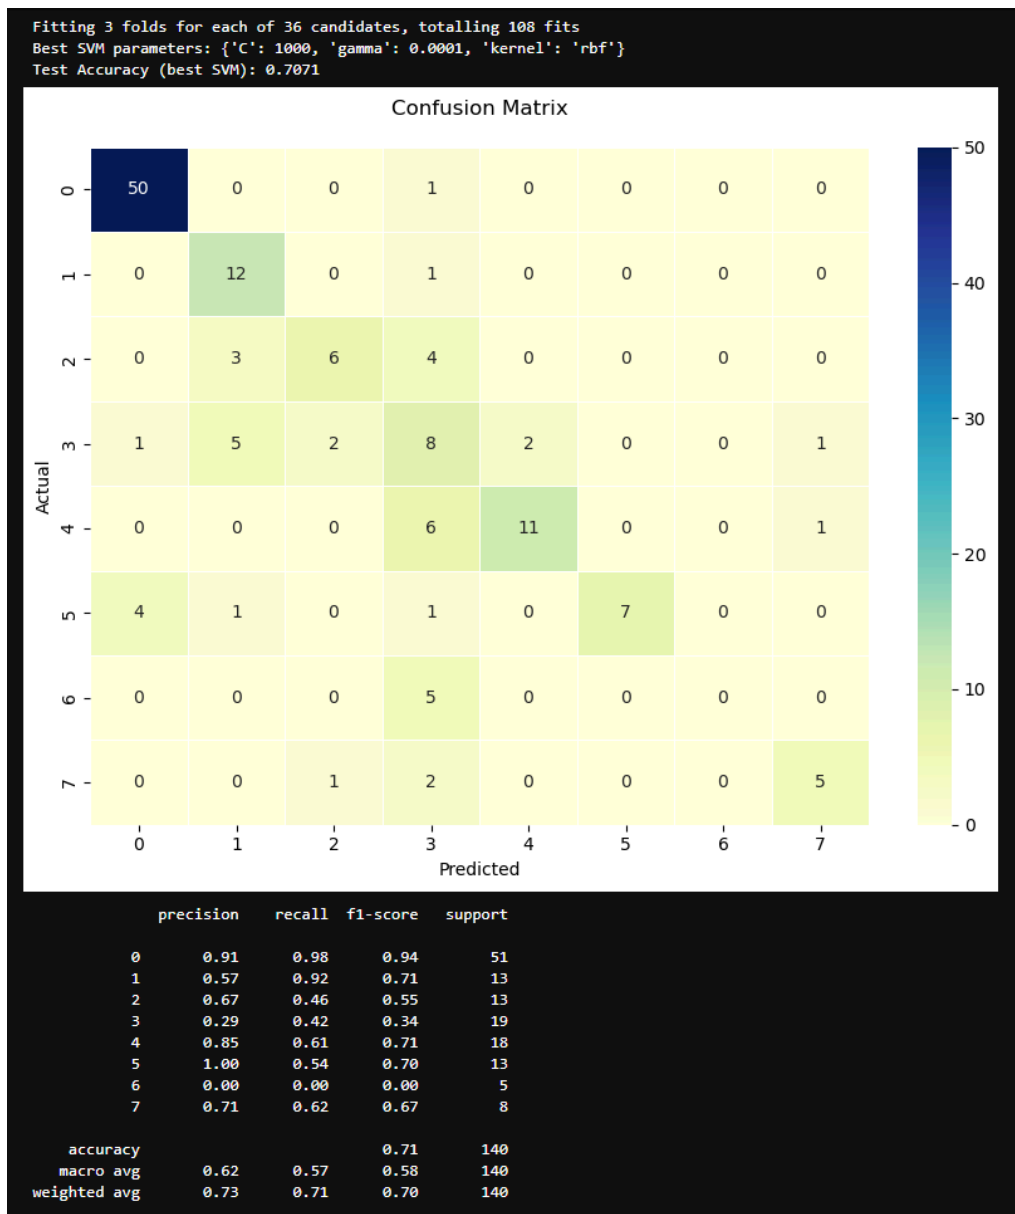
True: NotASolarEclipse  
Pred: 56to95percentPartialEclipse



# SIFT + BoVW + SVM/LR

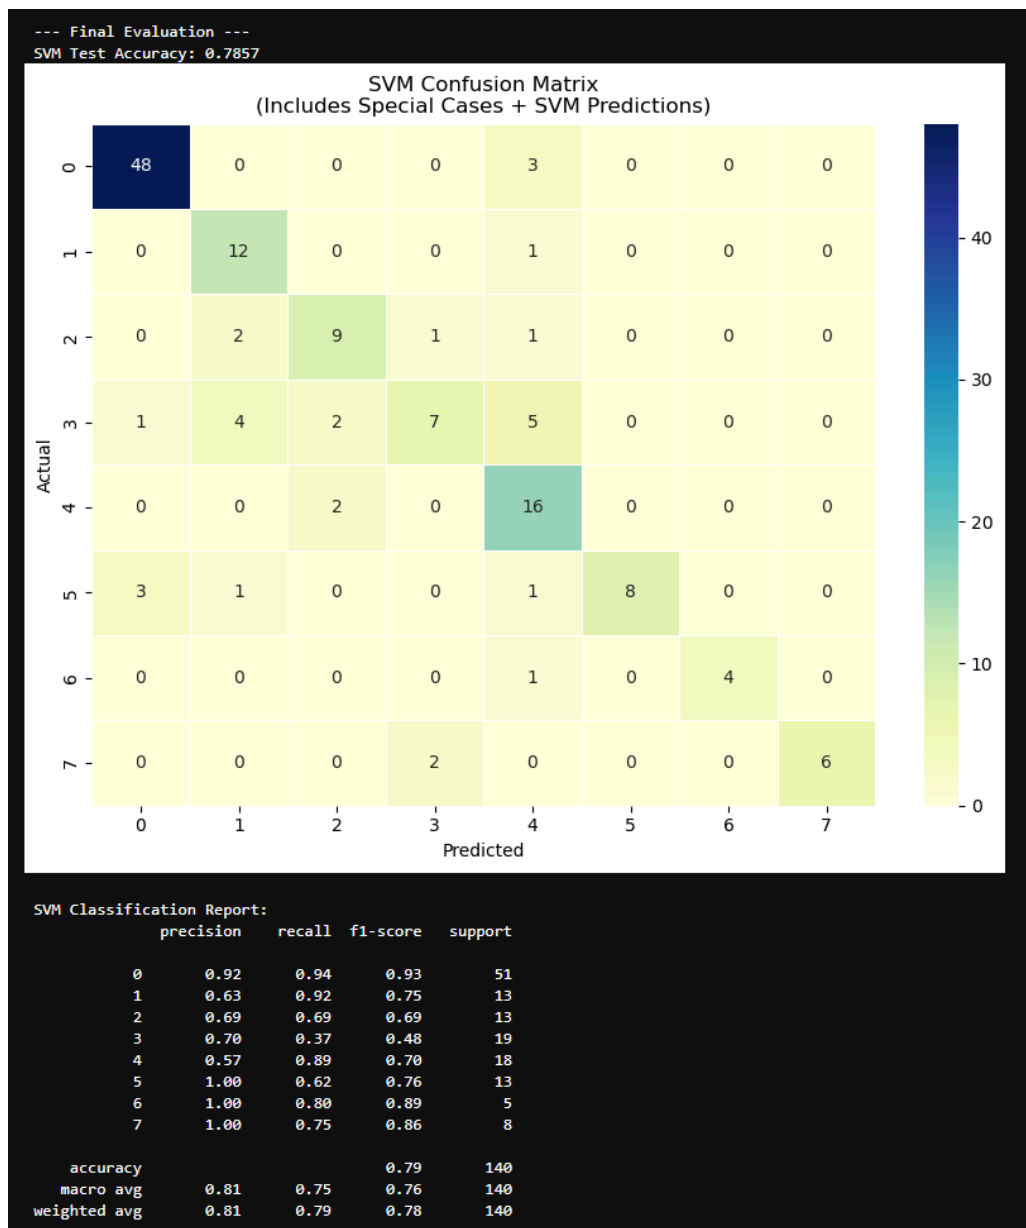
## First Pipeline

Our first attempt at using SIFT and BoVW did not use pre-classification of dark and flat images. Instead we went directly by trying to run all images through SIFT and BoVW and then through SVM and Logistic Regression. The following were our best results:



## Second Pipeline


No matter what parameters we changed we could not get a higher accuracy. So we decided to look at what images were mostly being misclassified based on the confusion matrix and we can see that 6(flats) had no correct predictions and 4(darks) could do better. We decided that we could classify these two beforehand since the key feature in each of these classes was simply the unbalanced histograms. Darks would have more of their peaks towards 0 and flats (brighter images) would have more of their peaks closer to 255. Using this information and intensities of these images, along with the dominant pixel value, we went ahead and tried classifying these images before running the rest of the images through SIFT, BoVW, and SVM/LR. With this we got the following results (best):





Logistic Regression Test Accuracy: 0.7643

Logistic Regression Confusion Matrix  
(Includes Special Cases + LR Predictions)



Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.92	0.92	0.92	51
1	0.67	0.92	0.77	13
2	0.73	0.62	0.67	13
3	0.60	0.32	0.41	19
4	0.55	0.89	0.68	18
5	0.80	0.62	0.70	13
6	1.00	0.80	0.89	5
7	0.86	0.75	0.80	8
accuracy			0.76	140
macro avg	0.77	0.73	0.73	140
weighted avg	0.78	0.76	0.76	140

## Overview of what was tried (for SIFT/BoVW/SVM/LR) and what the outcomes were:

- As stated, first we did **not try dark and flat classifying before SIFT**. Everything was running through sift and the **best accuracy** we were able to get (with varying ALL parameters) was **about 70%**.
- **Now pre-classifying flats and darks** through statistical thresholding, bumped the accuracy to our **best accuracy** of **about 79%**.
- With **this approach** our accuracy **varied from 74% up to 79%**. We tried **sharpening** the images since SIFT looks for keypoints that come from corners and edges (for example) because some eclipse images had good edges and points that stuck out, but we actually got a **worse accuracy** even with pre-classification i was getting **around 62%**.
- We tried using a **different preprocessing method for dark/flat classification** than for sift but the accuracy went down to **66%**. I believe we had tried **no CLAHE for flat/dark classification** and **no Gaussian Blur for SIFT/BoVW**.
- Parameters we played with and added:
  - `sift = cv2.SIFT_create(contrastThreshold=0.03, edgeThreshold=5)`
  - **Image size**
  - **Threshold values**
    - `def classify_flat_or_dark(img, flat_std_thresh=8, dark_mean_thresh=25, dark_std_thresh=6, path=""):`
  - **Minikmeans** instead of regular KMeans
    - Lowered the accuracy
  - **Vocab size**
    - Best: 345
    - Larger: Lowered the accuracy
    - Smaller: Lowered the accuracy
  - Added **grid search**
    - **Improved by about 1%** as our parameters before were almost all the same except for gamma, but always good to have